

Asynchronous Messaging

with ZMQ

Julius Parulek


DSD SOFTWARE ENGINEERING 2
Scout (Ert, Everest, ..., FMU)
Bergen

Asynchronous Messaging

with ZMQ (for running a foodtruck)

Julius Parulek


DSD SOFTWARE ENGINEERING 2
Scout (Ert, Everest, ..., FMU)
Bergen



Distributive systems

- Complex system of independent compute resources that appear as a single coherent system to the user
- **Features**
 - scalability
 - fault tolerance
 - distributive computation
 - distributive storage
 - resilient and efficient processing


Background



ØMQ

- High-Performance Messaging Library
- Socket Abstractions
 - Support different protocols
- **Asynchronous** design
- Built-in **messaging patterns**
- Peer-to-Peer Architecture
 - **Brokerless**
- Language Agnostic
 - C++, Java, Rust, Ruby, .Net, **Python**, ...

Background



Background

Foodtruck

- (chatgpt) imagine a restaurant and a delivery van had a **baby that makes delicious food**
- (wiki) A food truck is a large motorized vehicle or trailer equipped to **store, transport, cook, prepare, serve**, and/or sell food



schematics



schematics



schematics



schematics



[illegible]

1. Order

4. Prepare

Request-Reply

Request-Reply

- Bidirectional communication
- Synchronous
- Simple



Request-Reply



Request-Reply



Request-Reply

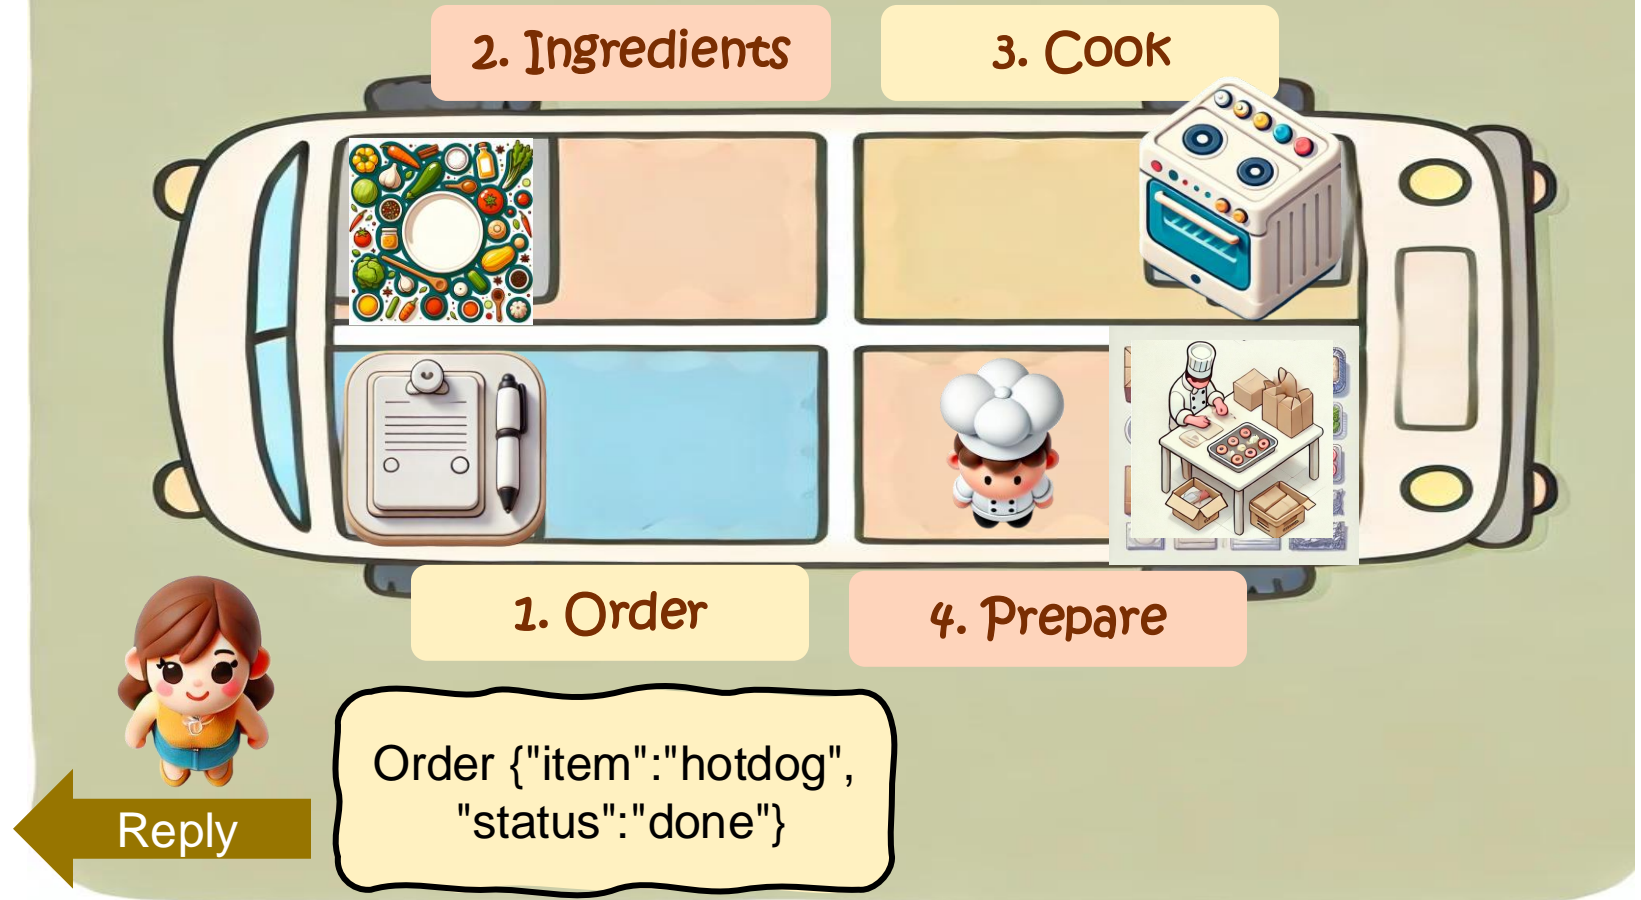


Request

Request-Reply



Request-Reply



Request-Reply



```
def customer():  
    context = zmq.Context()  
    socket = context.socket(zmq.REQ)  
    socket.connect("tcp://localhost:5555")
```

```
while True:  
    order = {"item": "hotdog"}  
    socket.send_json(order)  
    order = socket.recv_json()
```

Request-Reply

1. Cook

```
def foodtruck():  
    context = zmq.Context()  
    socket = context.socket(zmq.REP)  
    socket.bind("tcp://*:5555")
```

```
while True:  
    order = socket.recv_json()  
  
    print(f"Received order: {order}")  
    order = do_order(order)  
    order = ingredients(order)  
    order = cook(order)  
    order = prepare(order)  
  
    socket.send_json(order)
```



Request-Reply

So ...

- **Pros**

- Elegant & simple
- Stateful

- **Cons**

- Fully synced
- Not resilient
- Bad for multiple receivers
 - Socket blocked until replies comes



Exclusive Pair

Extension to Request-Reply

- Socket is **not** blocked

2. Ingredients



3. Cook



```
def customer():  
    context = zmq.Context()  
    socket = context.socket(zmq.PAIR)  
    socket.connect("tcp://localhost:5555")
```

1. C

```
def foodtruck():  
    context = zmq.Context()  
    socket = context.socket(zmq.PAIR)  
    socket.bind("tcp://*:5555")
```


What about efficiency?

- Step-1 - 20 secs
- Step-2 - 20 secs
- Step-3 - 180 secs
- Step-4 - 20 secs
- 1 transaction
 - 240 secs = 4 mins

Exclusive Pair



What about efficiency?

- Step-1 - 20 secs
- Step-2 - 20 secs
- Step-3 - 180 secs
- Step-4 - 20 secs
- 1 transaction
 - 240 secs = 4 mins
 - 16 customers in 1 hour

Exclusive Pair



More workers!



More workers with Req-Rep/Pair



More workers with Req-Rep/Pair

So ...

- **Order** wants to say something to **Ingredients**
 - **Ingredients** needs to stop & reply
 - **Ingredients** can be busy, **Order** needs to wait



More workers with Req-Rep/Pair

Req/Rep/Pair is not great for multiple connections and workers



More workers

What we need?

- **Asynchronous** messaging & execution
 - **Push / Pull**
 - **Publish / Subscribe**
 - **Router / Dealer**



Push-Pull

Push-Pull

- Distribute jobs across many workers
- **Scatter-Gather**
- Load balancing
- Work parallelization
- Messages **are not duplicated**



Push-Pull

```
{"item": "hotdog",  
"do_stage": "ing"}
```

Master
queue

2. Ingredients

3. Cook

So what about a single
master queue
managed by
order stage ?

- All can pull
- Flexible workers,
each one can do all
4 jobs

1. Order

4. Prepare

Order {"item": "hotdog"}



Push-Pull

```
{"item": "hotdog",  
"do_stage": "ing"}
```

Master
queue

2. Ingredients

3. Cook

So what about a single
master queue
managed by
order stage ?

- All can pull
- Flexible workers,
each one can do all
4 jobs

1. Order

4. Prepare

Order {"item": "hotdog"}



Push-Pull

`{"item": "ice-cream",
"do_stage": "ing"}` `{"item": "hotodg",
"do_stage": "ing"}`

Master
queue

2. Ingredients

3. Cook

So what about a single
master queue
managed by
order stage ?

- All can pull
- Flexible workers,
each one can do all
4 jobs

1. Order

4. Prepare



```
{"item": "ice-cream", "do_stage": "ing"}  
{"item": "hotodg", "do_stage": "ing"}
```

Master
queue

Push-Pull

2. Ingredients

3. Cook

So what about a single
master queue
managed by
order stage ?

- All can pull
- Flexible workers,
each one can do all
4 jobs

1. Order

4. Prepare



`{"item": "ice-cream", "do_stage": "ing"}`
`{"item": "hotodg", "do_stage": "ing"}`

Master
queue

Push-Pull

2. Ingredients

3. Cook

So what about a single
master queue
managed by
order stage ?

- All can pull
- **Flexible** workers,
each one can do all
4 jobs

```
while True:
    order = await master_queue.recv_json()
    worker_name = order["do_stage"]
    while not free_stage(worker_name):
        await asyncio.sleep(0.1)
    order = await do_work(worker_name, order)
    master_queue.send_json(order)
```

h-Pull

2. Ingre

Tight coupling
between stages!

Push-Pull
for jobs of the
same type

one can do all
4 jobs

```
while True:
    order = await master_queue.recv_json()
    worker_name = order["do_stage"]
    while not free_stage(worker_name):
        await asyncio.sleep(0.1)
    order = await do_work(worker_name, order)
    master_queue.send_json(order)
```


Cook

Push-Pull

2. Ingredients

3. Cook

So what about a message **queue** between every **stage** ?

- Ex. **Ingredients** queue
 - **Order stage** does push
 - **Ingredients stage** does pull

Prep

1. Order

4. Prepare

Order



Cook

Push-Pull

2. Ingredients

3. Cook

Ing.

Prep

1. Order

4. Prepare

Order



Cook

Push-Pull

2. Ingredients

3. Cook

Ing.

Prep

1. Order

4. Prepare

Order

Order {"item":"hotdog"}

Request



Cook

Push-Pull

2. Ingredients

3. Cook


Ing.

Prep

1. Order

4. Prepare

Order



Cook

Push-Pull

2. Ingredients

3. Cook

`{"item": "hotdog",
"order_id": 12,
"order": "ok"}`

Ing.

Prep

1. Order

4. Prepare

Order



Push-Pull

Cook

2. Ingredients

3. Cook

`{"item": "hotdog",
"order_id": 12,
"order": "ok"}`

Ing.

Prep

1. Order

4. Prepare

Order



Cook

Push-Pull

2. Ingredients

3. Cook

Ing.

Prep

1. Order

4. Prepare

Order



Push-Pull

```
{ "item": "hotdog",  
  "order_id": 12,  
  "order": "ok",  
  "ing": "ok" }
```

Cook

2. Ingredients

3. Cook

Ing.

Prep

1. Order

4. Prepare

Order



```
{ "item": "hotdog",  
  "order_id": 12,  
  "order": "ok",  
  "ing": "ok" }
```

Cook

Push-Pull

2. Ingredients

3. Cook

Ing.

Prep

1. Order

4. Prepare

Order



Cook

Push-Pull

2. Ingredients

3. Cook


Ing.

Prep

1. Order

4. Prepare

Order



Cook

Push-Pull

2. Ingredients

3. Cook

Ing.

```
{"item": "hotdog",  
  "order_id": 12,  
  "order": "ok",  
  "ing": "ok", "cook": "ok"}
```

Prep

1. Order

4. Prepare

Order



Cook

Push-Pull

2. Ingredients

3. Cook

Ing.

Prep

1. Order

4. Prepare

Order



Cook

Push-Pull

2. Ingredients

3. Cook

Ing.

Prep

1. Order

4. Prepare



```
{"item": "hotdog",  
  "order_id": 12,  
  "order": "ok", "ing": "ok",  
  "cook": "ok", "prep": "ok"}
```

Order

Cook

Push-Pull

2. Ingredients

3. Cook

Ing.

Prep

1. Order

4. Prepare



Reply

```
Order {"item": "hotdog",  
      "status": "complete"}
```

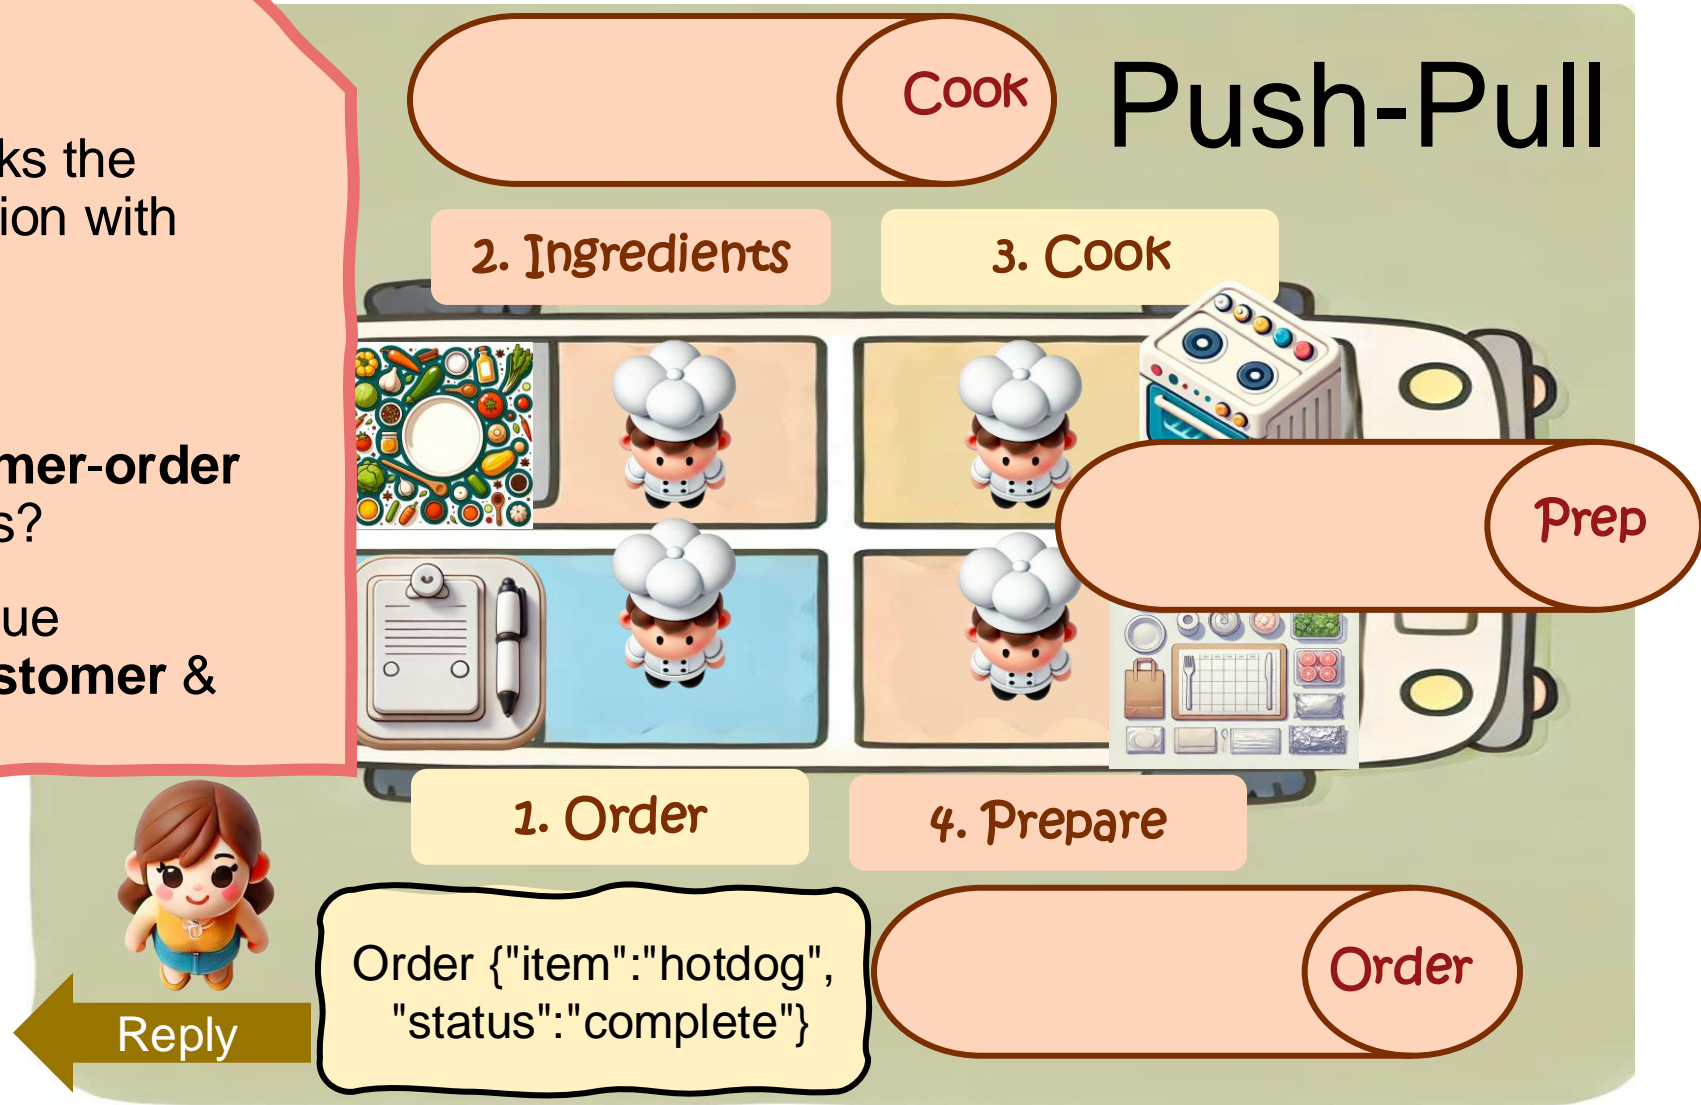
Order

But ...

- req-rep blocks the communication with customers

Fix ...

- More **customer-order** req-rep ports?
- Another queue between **customer & order**?



Push-Pull

Cook

Ing.

2. Ingredients

3. Cook

`{"item": "hotdog"}`

Order

Prep

1. Order

4. Prepare

Cust.



Push-Pull

Cook

```
{  
  "item": "hotdog",  
  "order_id": 12,  
  "order": "ok"  
}
```

Ing.

2. Ingredients

3. Cook

Order

Prep

1. Order

4. Prepare

Cust.



Push-Pull

Cook

```
{"item": "hotodog",  
"order_id": 12,  
"order": "ok", "ing": "ok"}
```

Ing.

2. Ingredients

3. Cook

Prep

```
{"item": "ice-cream"}
```

Order

1. Order

4. Prepare

Cust.



Cook

Push-Pull

```
{ "item": "ice-cream",  
  "order_id": 13,  
  "order": "ok" }
```

Ing.

2. Ingredients

3. Cook

Order

Prep

1. Order

4. Prepare

```
{ "item": "hotdog",  
  "order_id": 12,  
  "order": "ok", "ing": "ok",  
  "cook": "ok", "prep": "ok" }
```

Cust.




```
async def worker(worker_name):
    context = zmq.asyncio.Context()

    pull_socket = context.socket(zmq.PULL)
    pull_socket.connect(f"tcp://localhost:{ports['prep']}")

    push_socket = context.socket(zmq.PUSH)
    push_socket.bind(f"tcp://*:{ports['cust']}")

    while True:
        order = await pull_socket.recv_json()
        order = await do_work(worker_name, order)
        push_socket.send_json(order)
```

Push-Pull

3. Cook

Prep

1. Order

4. Prepare

Cust.



Push-Pull

Cook

2. Ingredients

3. Cook

Prep

Order

1. Order

4. Prepare

Cust.

```
{ "item": "ice-cream",  
  "order_id": 13,  
  "order": "ok" }
```

Ing.



Push-Pull

~~{"item": "ice-cream",
"order_id": 13,
"order": "ok", "ing": "ok"}~~

Cook

2. Ingredients

3. Cook

{"item": "ice-cream",
"order_id": 13,
"order": "ok"}

Skip!

Order

Prep

1. Order

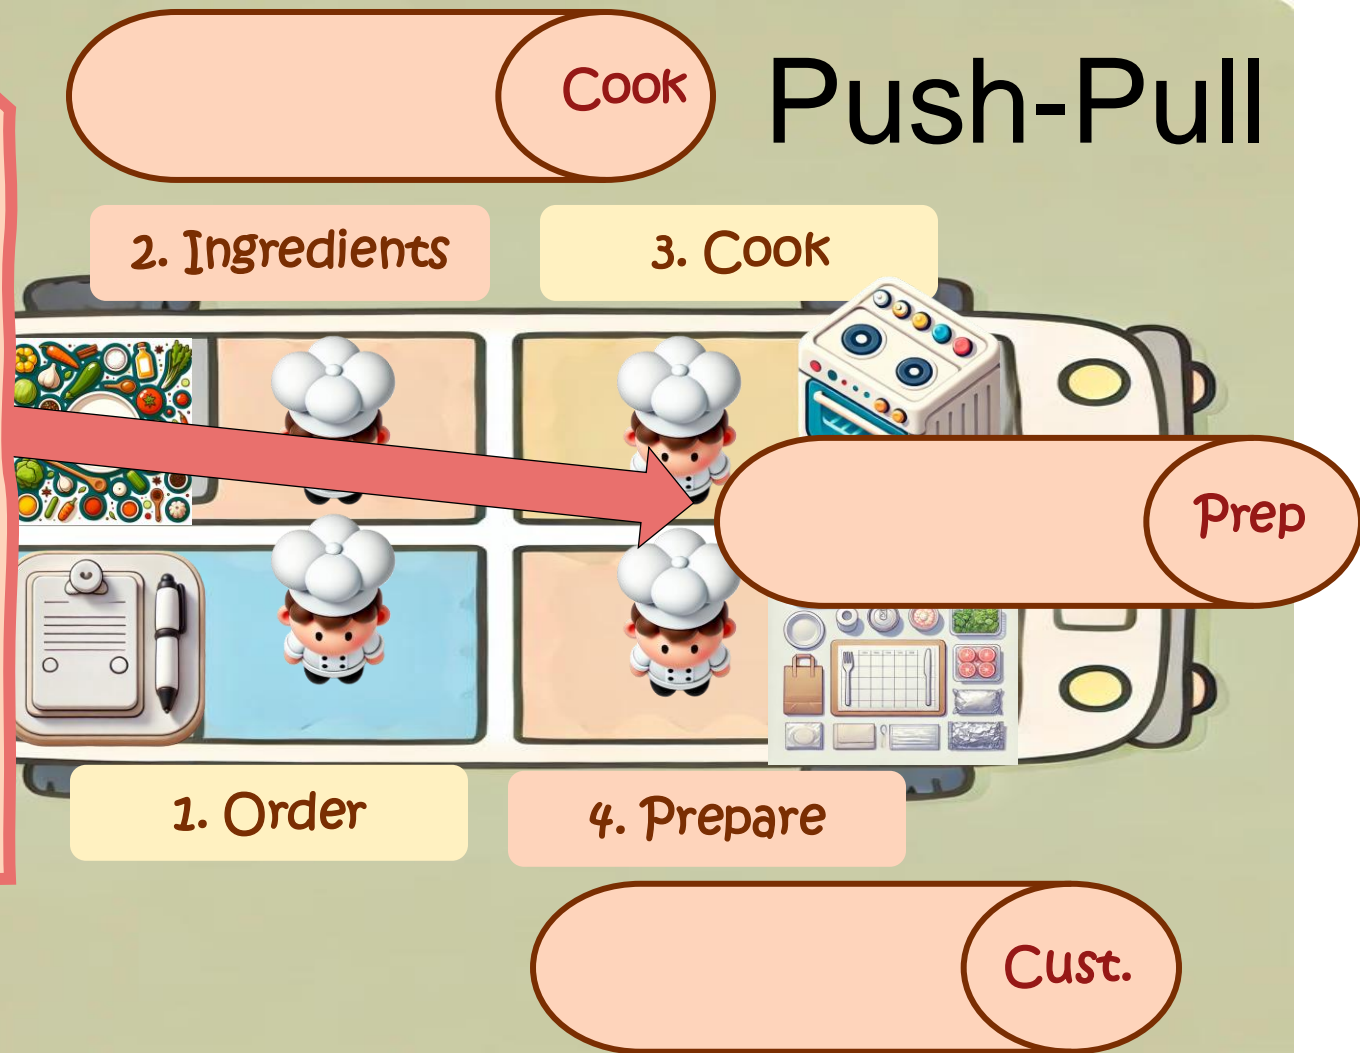
4. Prepare

Cust.



So ...

- Having **fixed** queues makes it difficult to skip **stages**
- Can be solved by connecting each stage as contributor (**push**) to each queue



Push-Pull

So ...

- **Pros**

- Elegant & simple
- Stateless
- Scalable
- **Parallelization**
- Fully async

- **Cons**

- Foodtruck **fails** since a stage can fail and the **customer** is not informed
- Integrating a new or skipping stages (ice-cream)
- Customer's direct access to worker
- **Jobs of the same type**

Ingredients

3. Cook



Publish-Subscribe

Features

- **One** sender, **multiple** receivers
- Subscribers can filter out messages – **topics**
- A new actor - **chef**



all:

chef:

Publish-Subscribe



all:

Publish-Subscribe

2. Ingredients

3. Cook

1. Order

4. Prepare



[cust, order, ing, cook, prep]

chef:

topic: order

{"item":"hotdog"}



[chef]

all:

topic: **order**
{ "item": "hotdog" }

Publish-Subscribe

2. Ingredients

3. Cook

1. Order

4. Prepare



[cust, order, ing, cook, prep]

chef:

[chef]

all:

topic: **order**
{ "item": "hotdog" }

Publish-Subscribe

2. Ingredients

3. Cook

1. Order

4. Prepare

[cust, order, ing, cook, prep]

chef:

[chef]



all:

topic: **order**

`{"item": "hotdog"}`

[cust, order, ing, cook, prep]

chef:

[chef]

Publish-Subscribe



unit:

Publish-Subscribe

2. Ingredients

3. Cook

1. Order

4. Prepare



[cust, order, ing, cook, prep]

chef:

topic: ingredients
{ "item": "hotdog",
 "order_id": 12,
 "order": "ok" }



all:

topic:

ingredients

```
{"item": "hotdog",  
  "order_id": 12,  
  "order": "ok"}
```

[cust, order, ing, cook, prep]

chef:

topic:

cook

```
{"item": "hotdog",  
  "order_id": 12,  
  "order": "ok",  
  "ing": "ok"}
```

[chef]

Publish-Subscribe



all:

topic:

ingredients

```
{"item": "hotdog",  
"order_id": 12,  
"order": "ok"}
```

[cust, order, ing, cook, prep]

chef:

topic:

cook

```
{"item": "hotdog",  
"order_id": 12,  
"order": "ok",  
"ing": "ok"}
```

Publish-Subscribe

```
while True: # ingredients job  
    msg = await all_socket.recv_string()  
    _, json_data = msg.split(" ", 1)  
    order = json.loads(json_data)  
    order = await do_work("ingredients", order)  
    new_topic = "cook"  
    json_data = json.dumps(order)  
    await chef_socket.send_string(new_topic + " " + json_data)
```



1. Order

4. Prepare

all:

topic: **cook**
{ "item": "hotdog",
 "order_id": 12,
 "order": "ok",
 "ing": "ok" }

[cust, order, ing, cook, prep]

chef:

topic:
prepare
{ "item": "hotdog",
 "order_id": 12,
 "order": "ok",
 "ing": "ok",
 "cook": "ok" }

[chef]

Publish-Subscribe



all:

topic:

prepare

```
{"item": "hotdog",  
  "order_id": 12,  
  "order": "ok",  
  "ing": "ok",  
  "cook": "ok"}
```

[cust, order, ing, cook, prep]

chef:

topic: customer

```
{"item": "hotdog",  
  "order_id": 12,  
  "order": "ok",  
  "ing": "ok",  
  "cook": "ok",  
  "prep": "ok"}
```

[chef]

Publish-Subscribe



all:

topic: **customer**
{**"item":**"hotdog",
"status":"done"}

[cust, order, ing, cook, prep]

chef:

topic: customer
{**"item":**"hotdog",
"order_id":12,
"order":"ok",
"ing":"ok",
"cook":"ok",
"prep":"ok"}

Publish-Subscribe



all:

topic: ingredients
{ "item": "hotdog",
"order_id": 12 }

What-if?

[cust, order, ing, cook, prep]

chef:

topic: ingredients
{ "item": "hotdog",
"order_id": 12,
"issue": "no sausages" }

[chef]

Publish-Subscribe



all:

topic: **prepare**
{ "item": "sausage",
 "order_id": 20,
 "cmd": "buy 200" }

What-if?

[cust, order, ing, cook, prep]

chef:

[chef]

Publish-Subscribe



all:

topic: ingredients
{ "item": "ice-cream",
 "order_id": 13 }

What-if?

[cust, order, ing, cook, prep]

chef:

topic: cook
{ "item": "ice-cream",
 "order_id": 13 }

[chef]

Publish-Subscribe



all:

topic: **prepare**
{ "item": "ice-cream",
 "order_id": 13 }

What-if?

[cust, order, ing, cook, prep]

chef:

topic: **cook**
{ "item": "ice-cream",
 "order_id": 13 }

[chef]

Publish-Subscribe



Publish-Subscribe

So

- **Pros**

- Elegant & simple
- Stateless
- **Scalable**
- Loose coupling
- Fully async

- **Cons**

- Potential message loss if subscribers are offline
- Messages are **broadcasted**

Ingredients

3. Cook





Features

- **Router** can connect to multiple **dealers**
- Fits complex distributive systems
- Highly scalable!

Router-Dealer

2. Ingredients

3. Cook





Router-Dealer





identities:

[customer, order,
ingredients, cook,
prepare]

Router-Dealer





identities:

[customer, order,
ingredients, cook,
prepare]

customer
`{"item":"hotdog"}`



Router-Dealer





identities:

[customer, order,
ingredients, cook,
prepare]

order
`{"item": "hotdog"}`



Router-Dealer





identities:

[customer, order,
ingredients, cook,
prepare]

order

```
{"item": "hotdog",  
  "order_id": 12,  
  "order": "ok"}
```



Router-Dealer

2. Ingredients



3. Cook



1. Order

4. Prepare



identities:

[customer, order,
ingredients, cook,
prepare]

ingredients

```
{"item": "hotdog",  
  "order_id": 12,  
  "order": "ok"}
```



Router-Dealer





identities:

[customer, order,
ingredients, cook,
prepare]

ingredients

```
{"item": "hotdog",  
  "order_id": 12,  
  "order": "ok",  
  "ing": "ok"}
```



Router-Dealer





identities:

[customer, order,
ingredients, cook,
prepare]

cook

```
{"item": "hotdog",  
  "order_id": 12,  
  "order": "ok",  
  "ing": "ok"}
```



Router-Dealer





identities:

[customer, order,
ingredients, cook,
prepare]


cook

```
{"item": "hotdog",  
 "order_id": 12,  
 "ing": "ok",  
 "order": "ok"}
```

Router-Dealer

2. Ingredients

3. Cook



```
while True:  # cook  
    [_, msg] = await dealer_socket.recv_multipart()  
    order = json.loads(msg)  
    order = await do_work("cook", order)  
    separator = b"  
    await dealer_socket.send_multipart([separator, json.dumps(order).encode()])
```




identities:

[customer, order,
ingredients, cook,
prepare]

cook

```
{"item": "hotdog",  
  "order_id": 12,  
  "order": "ok",  
  "ing": "ok", "cook": "ok"}
```



Router-Dealer





identities:

[customer, order,
ingredients, cook,
prepare]

cook

```
{"item": "hotdog",  
"order_id": 12,  
"order": "ok",  
"ing": "ok", "cook": "ok"}
```

Router-Dealer

2. Ingredients



3. Cook



```
while True: # router  
    [sender_id, _, msg] = await router_socket.recv_multipart()  
    sender_name = sender_id.decode() # cook  
    receiver_name = next_step[sender_name] # prep  
    await router_socket.send_multipart([receiver_name.encode(), b"", msg])
```



identities:

[customer, order,
ingredients, cook,
prepare]

prepare

```
{"item": "hotdog",  
"order_id": 12,  
"order": "ok", "ing": "ok",  
"cook": "ok"}
```



Router-Dealer





identities:

[customer, order,
ingredients, cook,
prepare]

prepare

```
{ "item": "hotdog",  
  "order_id": 12, "order": "ok",  
  "ing": "ok", "cook": "ok",  
  "prep": "ok" }
```



Router-Dealer





identities:

[customer, order,
ingredients, cook,
prepare]

customer

```
{"item": "hotdog",  
  "status": "done"}
```



Router-Dealer





identities:

[customer, order,
ingredients, cook,
prepare]

Router-Dealer





identities:

[customer, order,
ingredients, cook,
prepare]

Handles ...

- redundant & new stages
- missing ingredients & other issues

Can be solved by the
router!

Router-Dealer



So ...



- **Pros**

- Scalable
- Stateful
- Universal
- Fully async

- **Cons**

- More complex to implement and manage
- Weak coupling (identities)
- Careful design

Router-Dealer

Ingredients

3. Cook



1. Order

4. Prepare

Protocols in ØMQ

- **tcp://**
 - Across network communications
- **udp://**
 - Across network communications, receiver side no guaranteed
- **ipc://**
 - Inter-process communications
- **inproc://**
 - Same process, multi-thread communication

ents

3. Cook

```
radio = context.socket(zmq.RADIO)
radio.bind("udp://*:5555")

dish = context.socket(zmq.DISH)
dish.connect("udp://localhost:5555")
```

```
req = context.socket(zmq.REQ)
req.connect("ipc:///tmp/zmq-ipc")

rep = context.socket(zmq.REP)
rep.bind("ipc:///tmp/zmq-ipc")
```

More about ØMQ

- **Messages**
 - Binary objects
 - Any serializable object
- **Security**
 - Can be paired with CurveZMQ to provide encryption
- **Low level socket monitoring**
 - Validate various events occurring on the socket

nts

3. Cook



Summary

- **Real world systems** often use combination of these pattern
 - Router-Dealer with Publisher-Subscriber for scalable and event driven systems
- Consider system **requirements**
 - Is scalability required?
 - Loose or tight coupling?
 - Message deliver guarantees?



Further Reading

Books

- *ZeroMQ: Messaging for Many Applications* (Pieter Hintjens)
- *Enterprise Integration Patterns: Designing, Building, and Deploying Messaging Solutions* (Hohpe Gregor and Woolf Bobby)
- *Building Microservices: Designing Fine-Grained Systems* (Sam Newman)

Web

- <http://wiki.zeromq.org/>

Examples from presentation

- <https://github.com/xjules/edc2024-messaging>

Further Reading

Books

- *ZeroMQ: Messaging for Many Applications* (Noah Lerner)
- *Enterprise Integration Patterns: Designing Messaging Systems* (Gregor Hohpe)
- *Building a Reactive Microservices Architecture* (Sam Newman)

Web

- <http://wiki.zeromq.org/>

Examples from presentation

- <https://github.com/xjules/edc2024-messaging>

Thank you for your attention!